

Emotion Recognition

Arashdeep Mehroke

Mohamed Hasan

Dec 5, 2024

1 Problem Statement

Recognizing human emotions is an essential aspect of understanding and improving human-computer interaction, providing insights into behavioral analysis, mental health monitoring, and customer experience optimization. The **Ryerson Audio-Visual Database of Emotional Speech and Song (RAVDESS)** offers a rich collection of labeled audio data that can be leveraged to develop robust machine learning models for emotion classification.

Despite the availability of advanced feature extraction techniques such as YAMNet embeddings and Librosa-based handcrafted features, achieving high accuracy in emotion recognition remains a challenge due to:

1. **High Dimensionality of Features:** Audio features like YAMNet embeddings are inherently high-dimensional, complicating the training process for traditional models.
2. **Imbalanced Datasets:** Emotional classes often exhibit uneven distributions, leading to biased model performance.
3. **Complexity of Audio Signals:** Variations in pitch, tone, and intensity across speakers make it difficult to accurately classify emotions.
4. **Computational Overheads:** Extracting, visualizing, and analyzing audio features for large datasets requires efficient preprocessing and dimensionality reduction.

This project aims to address these challenges by:

1. Employing state-of-the-art feature extraction techniques (e.g., YAMNet, Librosa).
2. Balancing the dataset using synthetic techniques like SMOTE.
3. Experimenting with various machine learning and deep learning models to identify optimal solutions.
4. Visualizing high-dimensional feature spaces with dimensionality reduction tools like t-SNE and UMAP for better interpretability.

The goal is to build a robust, scalable system capable of accurately classifying human emotions across multiple categories.

2 Data Source

The dataset used in this project is the **Ryerson Audio-Visual Database of Emotional Speech and Song (RAVDESS)**. It consists of 1440 audio files (16-bit, 48kHz .wav) featuring 24 professional actors (12 female, 12 male) vocalizing two statements in a neutral North American accent.

2.1 Emotions and Features

The dataset includes 8 emotions: neutral, calm, happy, sad, angry, fearful, surprised, and disgusted, with normal and strong intensities (except neutral). Each statement is repeated twice, resulting in 60 trials per actor.

2.2 Filename Structure

Files are named using a 7-part numerical identifier (e.g., 03-01-06-01-02-01-12.wav) indicating:

- **Modality:** Audio-only (03).
- **Vocal Channel:** Speech (01).
- **Emotion:** Fearful (06).
- **Intensity:** Normal (01).
- **Statement:** Dogs (02).
- **Repetition:** First (01).
- **Actor:** Actor 12 (female).

3 Methodology

3.1 Data Preprocessing

This study uses audio datasets derived from the RAVDESS dataset, which include features extracted using Librosa and YAMNet libraries. Below is a summary of the engineered datasets generated:

- **librosa_balanced.csv:** 45 features, 7119 samples.
- **librosa_extracted_features.csv:** 45 features, 6439 samples.
- **yamnet_balanced.csv:** 1025 features, 7119 samples.
- **yamnet_extracted_features.csv:** 1025 features, 6440 samples.

Feature Exploration

Figure 1 illustrates the distribution of labels in the dataset. Most emotion classes are balanced, but the classes "Neutral" and "Surprise" are underrepresented. For this reason, Synthetic Minority Oversampling Technique (SMOTE) is used to address class imbalances.

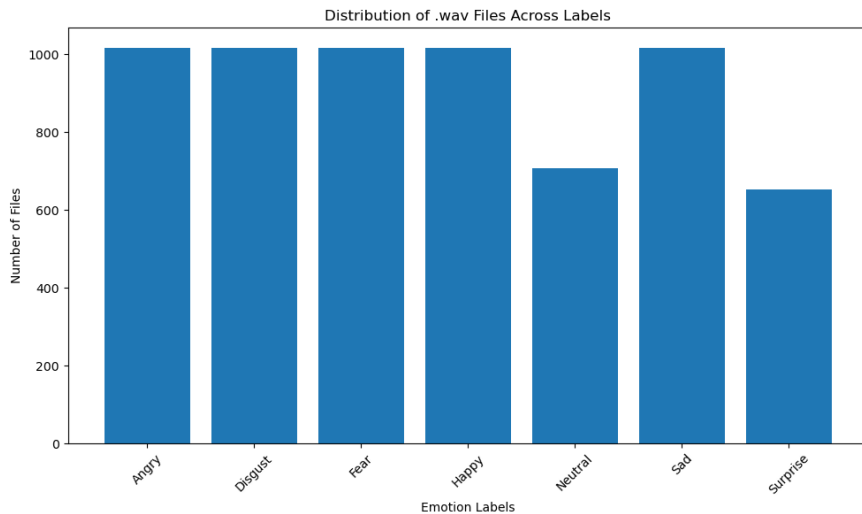


Figure 1: Distribution of emotion labels across audio samples.

Feature distributions across the dataset were analyzed. Figure 2 illustrates the distributions of key extracted features. The distributions reveal a variety of patterns in audio characteristics, which are useful for differentiating emotions.

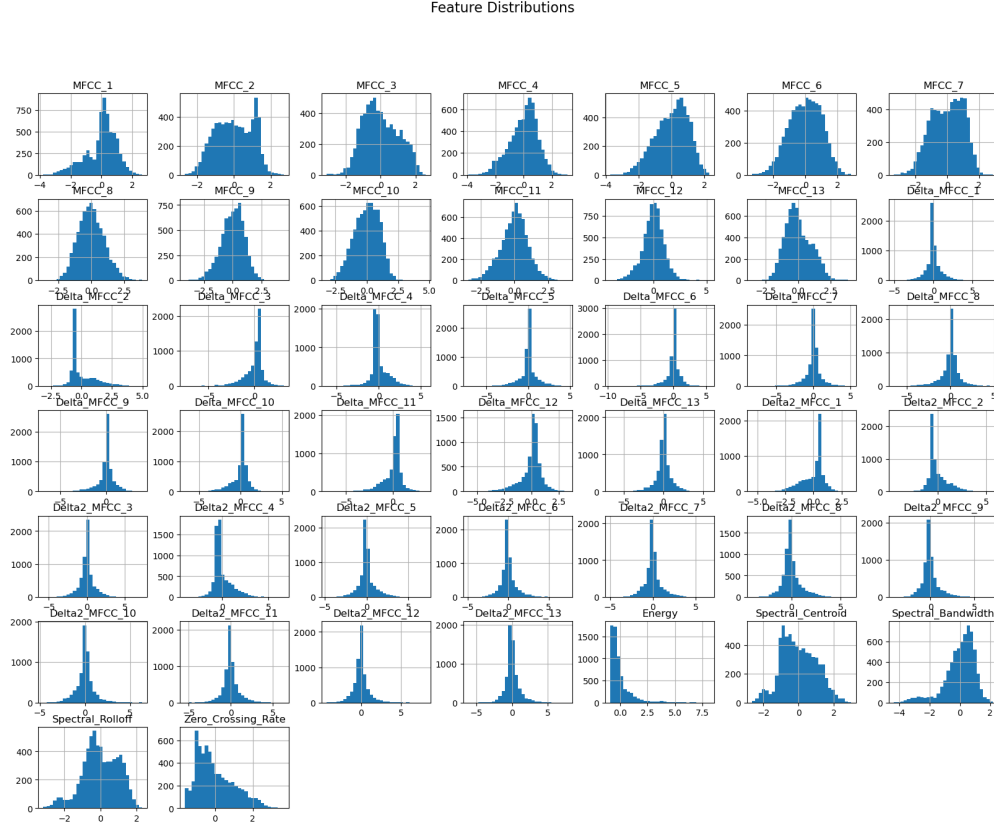


Figure 2: Feature distributions for extracted audio features, including MFCCs, spectral centroid, zero-crossing rate, and energy.

Feature Analysis

Figures 3 and 4 provide insights into the feature distributions and variability for the extracted features.

- **MFCCs:** The Mel-Frequency Cepstral Coefficients (MFCCs) capture timbral information. Peaks and valleys in the MFCC values are indicative of variations in speech timbre across emotions.
- **Spectral Centroid:** Represents the "brightness" of the audio. High centroid values suggest a higher frequency content, while lower values are linked to softer tones.
- **Zero-Crossing Rate (ZCR):** Indicates the noisiness or percussiveness of the signal. Higher ZCR values correspond to noisier audio signals, such as the "Surprise" class.
- **Energy:** Captures the loudness or intensity of the audio signal. Emotions like "Angry" and "Happy" show higher energy levels compared to "Sad" and "Neutral."

Dimensionality Reduction

Dimensionality reduction techniques such as PCA and t-SNE were applied to visualize the feature space. Figure 5 demonstrates clustering of the dataset using PCA. Despite overlapping clusters, PCA highlights trends in feature separability.




Figure 3: Mean values of MFCCs, Spectral Centroid, Zero-Crossing Rate, and Energy across emotions.

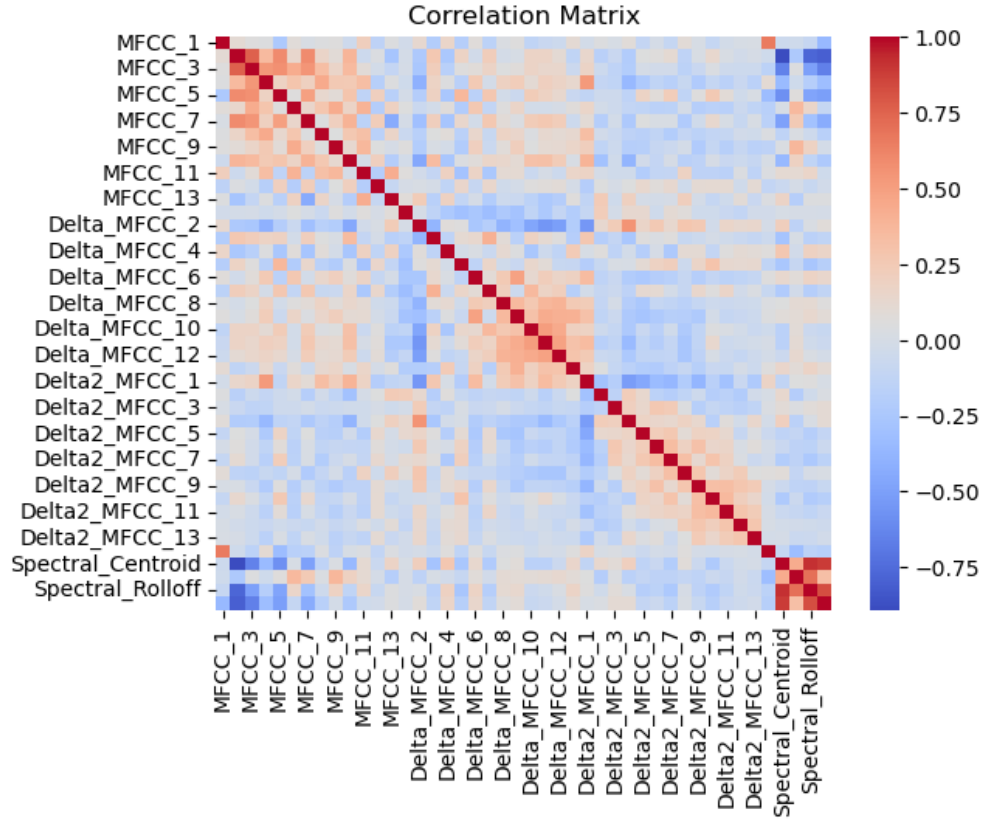


Figure 4: Distribution of feature values across all extracted features.




Figure 5: Clusters visualized using PCA. The clustering indicates some separation between emotion classes, though significant overlap is observed.

To address class imbalances, SMOTE (Synthetic Minority Oversampling Technique) was applied, and t-SNE was used to assess the improved separability of the dataset. Figure 6 highlights how t-SNE captures the non-linear relationships among features, improving the clustering of emotions.

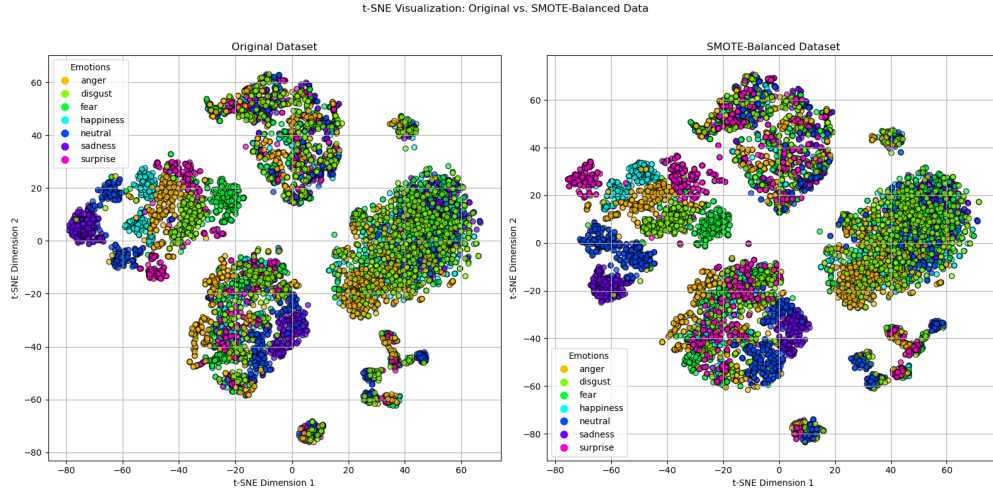


Figure 6: t-SNE visualization of original and SMOTE-balanced datasets. Non-linear feature relationships are effectively captured, enabling better emotion classification.

Model Selection

To classify emotions from the extracted audio features, four machine learning models were employed: K-Nearest Neighbors (KNN), Support Vector Machine (SVC), Random Forest, and a Multi-Layer Perceptron (MLP). Grid search was used to optimize the hyperparameters for KNN, SVC, and Random Forest, while MLP was manually tuned to align with the complexity of the problem.

- **K-Nearest Neighbors (KNN):** KNN was chosen for its simplicity and interpretability. It performs

well in scenarios where emotion classes have non-linear boundaries, as it classifies based on the proximity of similar feature patterns. KNN serves as an effective baseline for understanding how well the extracted features group emotions in the feature space.

- **Support Vector Machine (SVC):** SVC is particularly suited for high-dimensional datasets like YAMNet features, as it seeks to find the optimal margin between classes. Its ability to use kernel functions allows it to capture both linear and non-linear decision boundaries, making it a strong candidate for emotion classification where subtle differences in feature values distinguish emotions such as "Happy" and "Fearful."
- **Random Forest:** Random Forest was selected for its robustness and ability to handle heterogeneous data, such as a mix of spectral and temporal features. It combines decision trees to handle complex interactions among features, which is crucial in capturing the intricate variations in speech and audio characteristics associated with different emotions.
- **Multi-Layer Perceptron (MLP):** MLP was chosen for its ability to model complex, non-linear relationships in the data. Emotions often involve subtle, overlapping patterns in features like energy, spectral centroid, and MFCCs, which MLP can effectively learn. Its hierarchical feature learning capability makes it particularly well-suited for tasks like emotion recognition, though it requires careful tuning to prevent overfitting.

Hyperparameter Tuning

Grid search was applied to systematically explore the hyperparameter space for KNN, SVC, and Random Forest. For KNN, the number of neighbors and distance metrics were tuned; for SVC, the kernel type, regularization parameter (C), and kernel coefficient ($gamma$) were optimized; and for Random Forest, the number of trees ($n_estimators$), maximum tree depth (max_depth), and minimum samples per split ($min_samples_split$) were adjusted. MLP was tuned manually based on its learning rate, number of hidden layers, and neurons per layer.

Model Evaluation

All models were evaluated using metrics relevant to the emotion classification task: accuracy, precision, recall, and F1-score. These metrics were chosen to balance the trade-offs between false positives and false negatives, which are critical for applications like affective computing. Cross-validation was employed to ensure robust evaluation, particularly given the imbalanced nature of certain emotion classes. The results of these evaluations are discussed in detail in the subsequent sections.

Evaluation and Final Results

The performance of the selected models (KNN, SVC, Random Forest, and MLP) was evaluated on multiple datasets using accuracy, precision, recall, and F1-score as metrics. The results are summarized in Table 1 and visualized in Figures 7 and 8.

Summary of Results

Table 1 shows the performance metrics of all models across the different datasets. Accuracy is presented in a standalone graph (Figure 7) to emphasize its importance in emotion classification tasks, while precision, recall, and F1-score are grouped together in Figure 8 for a comprehensive comparison.

Visual Analysis of Results

The model accuracies are depicted in Figure 7, highlighting that the *SVC* consistently outperformed other models on most datasets. The *RandomForestClassifier*, while robust, showed slightly lower accuracy compared to *SVC* and *MLP*.

Table 1: Performance Summary of Models Across Datasets

Dataset	Model	Accuracy	Precision	Recall	F1-Score
yamnet_extracted_features	MLPClassifier	0.655	0.653	0.655	0.653
yamnet_balanced	MLPClassifier	0.697	0.696	0.697	0.696
yamnet_balanced	RandomForestClassifier	0.642	0.636	0.642	0.634
yamnet_extracted_features	SVC	0.680	0.679	0.680	0.679
librosa_extracted_features	SVC	0.679	0.679	0.679	0.678
librosa_balanced	KNeighborsClassifier	0.642	0.634	0.642	0.633
yamnet_balanced	SVC	0.721	0.719	0.721	0.719
yamnet_balanced	KNeighborsClassifier	0.614	0.609	0.614	0.603
librosa_balanced	RandomForestClassifier	0.711	0.711	0.711	0.704
yamnet_extracted_features	KNeighborsClassifier	0.598	0.600	0.598	0.597
librosa_balanced	SVC	0.713	0.710	0.713	0.708
librosa_balanced	MLPClassifier	0.686	0.684	0.686	0.684
yamnet_extracted_features	RandomForestClassifier	0.615	0.616	0.615	0.611
librosa_extracted_features	KNeighborsClassifier	0.600	0.599	0.600	0.596
librosa_extracted_features	RandomForestClassifier	0.679	0.680	0.679	0.675
librosa_extracted_features	MLPClassifier	0.641	0.641	0.641	0.640




Figure 7: Model Performance: Accuracy across datasets and models.

Precision, recall, and F1-score are shown in Figure 8. These metrics provide insights into how well the models balance false positives and false negatives, which is critical in emotion classification where misclassifications can have significant impact.

Discussion

The evaluation results indicate that the *Support Vector Classifier (SVC)* outperformed other model, achieving the highest accuracy, particularly on the *yamnet_balanced* dataset with an accuracy of **72.1%**. This demonstrates SVC’s ability to handle high-dimensional feature spaces and capture complex decision boundaries effectively.

The superior performance of the *Support Vector Classifier (SVC)* can be attributed to its ability to effectively handle high-dimensional feature spaces and identify optimal decision boundaries between complex




Figure 8: Model Performance: Precision, Recall, and F1-Score across datasets and models.

and overlapping classes. With datasets such as *yamnet_balanced*, which contain a high number of features (1025 in total), SVC’s use of the kernel trick allows it to map data into higher-dimensional spaces where it becomes linearly separable. This capability is particularly beneficial for emotion classification, as emotional states like “happy” and “surprise” or “sad” and “neutral” often share overlapping feature distributions. Additionally, SVC’s robustness to outliers and class imbalance, combined with hyperparameter optimization (e.g., tuning the kernel type and regularization parameter), further contributed to its strong performance. These characteristics make SVC a highly effective model for the nuanced and complex task of emotion recognition from audio data.

Conclusion

The evaluation results highlight the strengths and limitations of the models tested for emotion classification based on extracted audio features. Audio classification, particularly emotion recognition, presents significant challenges due to the complex and overlapping nature of audio features, as well as the variability in how emotions are expressed across speakers and contexts. Key observations from this study include:

- **SVC Achieved the Highest Accuracy:** Across the datasets, *Support Vector Classifier (SVC)* consistently delivered the highest accuracy, showcasing its ability to handle high-dimensional feature spaces (e.g., YAMNet features) and complex decision boundaries. The optimal hyperparameter tuning using grid search further contributed to its superior performance.
- **MLP Demonstrated Balanced Performance:** The *Multi-Layer Perceptron (MLP)* showed competitive performance, especially on the *yamnet_balanced* dataset, where it achieved close to the highest scores across all metrics. MLP’s ability to learn hierarchical feature representations allowed it to handle the intricacies of audio features effectively. However, it requires more careful tuning and computational resources to avoid overfitting.
- **Random Forest’s Robustness:** *Random Forest* provided stable but slightly lower performance compared to SVC and MLP. Its robustness and ability to handle heterogeneous data were evident, but it struggled with the subtle, high-dimensional relationships present in audio features, which are critical for distinguishing nuanced emotional expressions.

- **KNN as a Baseline:** *K-Nearest Neighbors (KNN)* served as a baseline and performed reasonably well, demonstrating its simplicity and effectiveness for non-linear feature spaces. However, its sensitivity to class imbalances and computational inefficiency in high dimensions limited its overall performance.

Challenges in Audio Emotion Classification: Emotion classification from audio data is inherently challenging due to:

- **Feature Overlap:** Emotional states like "happy" and "surprise" or "sad" and "neutral" often share overlapping spectral and temporal features, making them difficult to distinguish.
- **High Dimensionality:** Datasets such as those with YAMNet features contain hundreds of features, necessitating robust models and hyperparameter optimization to avoid overfitting.
- **Class Imbalance:** Underrepresented classes like "surprise" require careful handling to ensure the model does not favor dominant classes.
- **Speaker Variability:** Differences in tone, pitch, and accent across speakers can add noise to the classification task, further complicating model training.

Overall Findings: The results indicate that SVC is the most effective model for emotion classification in this study, particularly for high-dimensional datasets such as those with YAMNet features. Its ability to generalize across diverse datasets makes it an ideal choice for applications where precision and recall are critical. MLP demonstrated promising results and could benefit from further tuning and advanced architectures to potentially surpass SVC in future experiments.

Future Work: To address the complexities of audio emotion classification, future efforts can focus on:

- Advanced feature engineering and dimensionality reduction to improve feature quality and reduce computational complexity.
- Exploring deeper neural architectures or pre-trained audio models such as transformers for enhanced representation learning.
- Using ensemble methods to combine the strengths of SVC, MLP, and Random Forest for improved accuracy.
- Incorporating techniques such as domain adaptation to handle speaker variability and make models more robust.

This study highlights the critical importance of model selection and optimization in tackling the challenges of audio emotion classification, emphasizing the need for rigorous evaluation and future innovations in this domain.

Citation

Livingstone SR, Russo FA (2018). The Ryerson Audio-Visual Database of Emotional Speech and Song (RAVDESS). *PLoS ONE*, 13(5): e0196391. <https://doi.org/10.1371/journal.pone.0196391>.

For other uses, attribute as: "*The RAVDESS by Livingstone & Russo, licensed under CC BY-NA-SC 4.0.*"